

# Algorithmic Collusion Detection

Matteo Courthoud\*

January 31, 2021

## Abstract

Reinforcement learning algorithms are gradually replacing humans in many decision making processes, such as pricing in high-frequency markets. Recent studies on algorithmic pricing have shown that algorithms can learn sophisticated grim-trigger strategies with the intent of keeping supracompetitive prices. One suggestion to detect algorithmic collusion is to look at the inputs of the dynamic strategies. In this paper, I show that this approach might not be sufficient since the algorithms can learn reward-punishment schemes that are fully independent from the rival's actions. Moreover, I show that the crucial ingredient in algorithmic collusion is synchronous learning. When one algorithm is unilaterally retrained, it learns more competitive strategies that exploit collusive behavior. Since this change in strategies happens only when algorithms are colluding, retraining can be used as an instrument to detect algorithmic collusion. Lastly, I show how one can get the same insights on collusive behavior by retraining the algorithms on bootstrapped samples of historical data.

**Keywords:** Artificial Intelligence, Collusion, Antitrust

**JEL Classification:** D21, D43, D83, L12, L13

---

\*University of Zürich, email: [matteo.courthoud@econ.uzh.ch](mailto:matteo.courthoud@econ.uzh.ch).

# 1 Introduction

Algorithms are slowly substituting human decision making in many business environments. While some applications such as playing chess or driving cars have attracted most of the attention of the press, algorithms are also more and more frequently used to make pricing decision. Since these algorithms need experience in the form of data to be trained, their applications mostly reside in high-frequency markets, such as gasoline markets (Assad et al., 2020), Amazon marketplace (Chen et al., 2016) or financial markets.

A recent strand of literature has highlighted the ability of these algorithms to learn complex dynamic collusive strategies, in the form of reward-punishment schemes. In particular, Calvano et al. (2020b) found that reinforcement learning algorithms are able to learn complex reward-punishment schemes with the intent of keeping supracompetitive prices. Most importantly, they learn these strategies autonomously, without any human intervention nudging them towards collusion.

Policymakers are currently discussing what can be done either to prevent these collusive behaviors or to detect and punish them. Ezrachi and Stucke (2017) and Ezrachi and Stucke (2019) suggest to create testing hubs for algorithms in order to classify them into potentially collusive and competitive. Another ex-ante solution, proposed by Harrington (2018), is to look at the algorithms' inputs and ban algorithms that are provided information that would allow them to collude, such as historical actions of their rivals. In contrast to this ex-ante solution, Calvano et al. (2020a) propose an interim approach: computer scientists should make algorithms more interpretable, so that humans could observe their behavior in real time and identify the rationales behind their decision making. Another potential approach is to intervene ex-post and sanction algorithms based on their observed behavior.

The first contribution of this paper is to show that, by solely inspecting the strategies of the algorithm, it is hard to draw conclusions on whether the algorithms are colluding. In particular, I show that pricing algorithms can learn reward-punishment schemes with the sole intent of keeping supracompetitive prices with strategies that are fully independent from competitors' actions. Indeed, algorithms learn to punish themselves for deviating from collusive prices and are hence able to obtain supracompetitive profits. While this extreme case is unrealistic, it underlines the difficulty of categorizing inputs into pro-collusive and pro-competitive.

Given that detecting collusion from the inspection of the algorithm strategies is particularly cumbersome, the second part of this paper proposes a model-free method to detect collusion strategies. My approach is based on the insight that the key aspect of algorithmic collusion is synchronous learning. I show that if one algorithm is retrained unilaterally, it learns to exploit its collusive competitors with a more competitive strategy.

The main drawback of this method is that it is infeasible in practice: requiring firms to retrain the algorithm is too expensive. Therefore, the last part of my work shows how to get the same insight from historical data, without the need to impose suboptimal actions on the firms employing pricing algorithms. Retraining algorithms on bootstraps of historical data allows to learn more competitive strategies to exploit the collusive behavior of their competitors. To the best of my knowledge, this constitutes the first attempt to build a model-free test for algorithmic collusion, using only historical data.

The paper is structured as follows. Section 2 contains a review of artificial intelligence algorithms in decision making. In particular, I cover recent advancements in computer science, with a focus on Q-learning algorithms and the specific formulation used in the simulations. In Section 3, I explore the extreme scenario in which a Q-learning pricing algorithm does not observe the competitor's price but still manages to adopt a reward-punishment strategy to keep supra-competitive prices. Section 4 focuses on collusion detection. I first show how unilaterally retraining one algorithm leads to more competitive strategies in case the algorithms were colluding. Then, I show how one can obtain the same insight from historical data. Section 5 concludes.

## 2 Q-Learning

Reinforcement learning algorithms are a class of learning algorithms that try to solve optimization problems where rewards are delayed over multiple time periods. Moreover, these rewards depend on the sequence of actions that the algorithm has to take over different time periods. These two characteristics make reinforcement learning fundamentally different from standard supervised learning problems. The algorithm objective is not simply to learn a reward function but also to take an optimal sequence of actions. The most important feature of reinforcement learning is indeed the dual role of the algorithm: prediction and optimization.

Q-learning is a popular reinforcement learning algorithm and constitutes the baseline model for the most successful advancements in Artificial Intelligence in the last decade (Igami, 2020). The most popular algorithms such as Bonanza or AlphaGo are based on the same principle, while using deep neural networks to provide a more flexible functional approximation of the policy function (Sutton and Barto, 2018).

One advantage of Q-learning algorithms is their interpretability, especially for what concerns the mapping from the algorithm to the policy and value functions. The policy function has a matrix representation that can be directly observed and interpreted. This makes it possible to understand not only the logic behind any decision of the algorithm at any point in time, but also to know what the algorithm would have done in any counterfactual scenario.

In this section, I first explore the general formulation of Q-learning algorithms. Since these algorithms have been developed to work in single-agent environments, I comment on their use in repeated games. Lastly, I analyze the baseline algorithm used in the simulations throughout the paper.

### 2.1 Single Agent Learning

Reinforcement learning algorithms try to solve complex dynamic optimization problems by adopting a model-free approach. This means that the algorithm is not provided any structure regarding the relationship between the state it observes, its own actions, and the payoffs it receives. The algorithm only learns through experience, associating states and actions with the payoffs they generate. Actions that bring higher payoffs in a state are preferred to actions that bring lower payoffs. Since the state can include past states or actions, reinforcement learning algorithms can learn complex dynamic strategies. The more complex the state and action space of the learning algorithm, the more complex the strategies it can learn.

The objective function of the learning algorithm is the expected discounted value of future payoffs

$$\mathbb{E} \left[ \sum_{t=0}^{\infty} \delta^t \pi_t \right]. \quad (1)$$

In many domains faced by computer scientists, payoffs have to be hard-coded together with the algorithm. For example, with self-driving cars, one has to establish what is the payoff of an accident, or the payoff of arriving late. Clearly, the decision of these payoffs directly affects the behavior of the algorithm. However, in some cases, the payoffs are directly provided by the environment. For example, in the setting analyzed in this paper, the algorithm sets prices and observes the profits coming from the sales of an item.

The main trade-off faced by reinforcement learning algorithms is the so-called *exploration-exploitation* trade-off. Since the algorithm is not provided with any model of the world, it only learns through experience. In order to learn different policies, the algorithm explores the action space by taking sub-optimal actions. However, since the objective of the algorithm is to maximize the expected discounted sum of future payoffs, at a certain point, the algorithm needs to shift from exploration to exploitation, i.e., it needs to start taking optimal actions, given the experience accumulated so far.

In each period, the algorithm observes the current state of the world,  $\mathbf{s}$  and takes an action  $a$  with the intent to maximize its objective function. We refer to the total discounted stream of future payoffs under optimal

actions as the value function. We can express the value function of algorithm  $i$  in state  $s$  recursively as

$$V_i(s) = \max_{a_i \in \mathcal{A}} \left\{ \pi_i(s, a) + \delta \mathbb{E}_{s'} [V_i(s') | s, a] \right\}. \quad (2)$$

Q-learning algorithms are based on a different representation of the value function, the action-specific value function, which is called the  $Q$  function. We can write the action-specific value function of algorithm  $i$  in state  $s$  when it takes action  $a_i$  as

$$Q_i(s, a_i) = \pi(s, a) + \delta \mathbb{E}_{s'} \left[ \max_{a'_i \in \mathcal{A}} Q_i(s', a'_i) | s, a \right]. \quad (3)$$

When the state space  $\mathcal{S}$  and the action space  $\mathcal{A}$  are finite, we can express the  $Q$  function as a  $|\mathcal{S}| \times |\mathcal{A}|$  matrix. Many of the advancements in reinforcement learning involve a functional representation of the  $Q$  function that can encompass more complex state or action spaces. The most successful function approximations involve deep neural networks.

The objective of the algorithm is to learn the  $Q$  function. Once the algorithm has learned the  $Q$  function, in each period it will take the optimal action  $a_i^* = \arg \max_{a_i \in \mathcal{A}} Q(s, a_i)$ . Learning and taking the optimal actions are the two main tasks of the algorithm, and the choice between the two behaviors constitutes the main trade-off of reinforcement learning: *exploration* versus *exploitation*.

*Exploitation.* Since the final objective of the algorithm is to maximize the expected discounted sum of payoffs, in the exploitation phase, the algorithm picks the best available action, given the experience accumulated so far. After a sufficient amount of exploration, in a stationary environment, the algorithm expected discounted sum of future payoffs is guaranteed to converge to a local optimum of the value function (Beggs, 2005). Exploration is needed in order to discover other better locally optimal policies.

*Exploration.* During the exploration phase, the objective of the algorithm is to explore the state-action space in order to discover new policies. Since exploration involves testing suboptimal actions, there is a trade-off between exploration and exploitation. More exploration implies lower short-term profits but can lead to the discovery of policies than bring higher long-term profits. Moreover, when reinforcement learning algorithms are deployed in a dynamic environment, exploration gives the algorithm the flexibility to adapt to changes in the environment.

There exists many different ways in which one algorithm can explore the state-action space and the simplest one is the  $\varepsilon$ -greedy model. In each period, the algorithm decides to explore with probability  $\varepsilon$  and to exploit with probability  $(1 - \varepsilon)$ . In the exploration phase, the algorithm chooses one action uniformly at random.

*Optimality Results.* In a Markov single agent environment, under mild conditions, it has been proven that a reinforcement learning algorithm learns locally optimal strategies, given that the exploration rate  $\varepsilon$  converges to zero as time goes to infinity (Sutton and Barto, 2018).

*Learning Speed.* Since learning is a noisy process, the  $Q$  matrix is only partially updated. In particular, given an action  $a_i^*$ , irrespectively of whether the action comes from exploration or exploitation, the update policy is

$$Q_i(s, a_i^*) = \alpha Q_i^{OLD}(s, a_i^*) + (1 - \alpha) Q_i^{NEW}(s, a_i^*) \quad (4)$$

where we refer to  $\alpha$  as the *learning rate*. A higher  $\alpha$  implies a faster but noisier learning. The policy function takes less time to converge but it is less likely to adopt better strategies.

## 2.2 Repeated Games

In our setting, multiple Q-learning algorithms play a repeated game in which they set per-period prices with the objective to maximize the expected discounted sum of profits. Most of the reinforcement learning literature in computer science focuses stationary environments. One common example is video games, where algorithms

receive the video feed as an input and have to pick the optimal actions in order to perform best in the game. Another common area of research in reinforcement learning is robotics, from logistics to self-driving cars. All these examples involve mostly stationary environments.

The behavior of reinforcement learning algorithms competing with each other in a repeated game is still under research and there exist no general result concerning their behavior. In particular, there is no result on whether algorithmic behavior will converge on collaborative behavior, depending on the context.

### 3 Blind Learning

Differently from humans, algorithms have strategies that are hard-coded and hence directly observable. We cannot observe whether humans decide to increase prices because of higher demand, higher input prices, or an agreement with competitors. This is indeed the reason why collusion per-se is not prohibited by law which, on the other hand, targets communication with the intent to collude. However, we can observe the decision-making process of the algorithm. Therefore, one might wonder whether it is possible to detect collusion by the inspection of the algorithms' strategies.

The main problem of detecting collusion from the inspection of algorithms' strategies comes from the fact that these strategies might be extremely complex and hard to understand for a human. In practice, most reinforcement learning algorithms do not rely on a matrix representation of the  $Q$  function but approximate it through deep neural networks which are known to be extremely difficult to interpret.

One approach that has been proposed by Harrington (2018) is to look at the inputs of the algorithm strategies. In particular, commenting on Calvano et al. (2020b), he suggests that one possible metric to flag collusive algorithms is to look at whether the algorithm prices are conditional on rivals' past prices.

*“In this simple setting, a pricing algorithm would be prohibited if it conditioned price on a rival firm’s past prices. AAs [Artificial Agents] would be allowed to set any price, low or high, but just not use pricing algorithms that could reward or punish a rival firm based on that firm’s past prices.”*

And he adds

*“I am not suggesting that, in practice, a pricing algorithm should be prohibited if it conditions on rival firms’ past prices. However, within the confines of this simple setting, that would be the proper definition of the set of prohibited pricing algorithms.”*

In this section, I am going to show that looking at the inputs of the algorithm strategies might not be sufficient to detect algorithmic collusion. In particular, I show that an algorithm whose strategy is based only on its own past action can still learn a reward-punishment scheme with the sole intent of keeping supra-competitive prices.

#### 3.1 Model

I adopt the baseline model from Calvano et al. (2020b). First of all, this choice allows direct comparison with their simulation results. Moreover, their setting is particularly simple and easy to interpret. Lastly, since many competition policy and law papers are based on their results, I ensure to directly speak to that literature.

Time is discrete and the horizon is infinite. At each point in time,  $n$  firms are active and compete in prices with differentiated products. Differently from Calvano et al. (2020b), the state of the game for each firm is its own history of pricing decisions up to  $k$  times periods in the past. We discretize the possible actions on a grid of dimension  $m$ :  $\{p_1, \dots, p_m\}$ . Therefore, the subjective state of firm  $n$  is represented by a vector  $\mathbf{s}_{i,t} = \{p_{i,t-1}, \dots, p_{i,t-k}\} \in S = \{p_1, \dots, p_m\}^k$ , where  $s_{i,t}$  represents the subjective state of firm  $i$  at time  $t$ . I will

refer to  $\mathbf{s}$  as a subjective state and  $S$  as the subjective state space. Firms maximize their total future discounted profits. The discount factor is  $\delta \in [0, 1)$ .

*Demand.* There is a continuum of consumers of unit mass. Consumer  $j$ 's utility from buying one unit of product  $i$  is given by

$$u_{j,i} = v_j - \mu p_i + \varepsilon_j \quad (5)$$

where  $v_j$  is the value of the product  $i$  for consumer  $j$ ,  $p_i$  is the price of product  $i$ ,  $\mu$  is the price elasticity and  $\varepsilon_j$  is the idiosyncratic shock preference of consumer  $j$  for product  $i$ . The random shocks  $\varepsilon_j$  are assumed to be independent and type 1 extreme value distributed so that the resulting demand function has the logit form. For example, the demand of product  $i$  is:

$$q_i(\mathbf{p}) = \frac{e^{-\mu p_i}}{e^{-\mu p_i} + \sum_{-i} e^{-\mu p_{-i}}}. \quad (6)$$

The model has a unique Nash Equilibrium which we refer to as the competitive outcome.

*Exploration/Exploitation.* We use a  $\varepsilon$ -greedy exploration method as in Calvano et al. (2020b). In each period, each algorithm has a probability  $\varepsilon_t$  of exploring and a probability  $1 - \varepsilon_t$  of exploiting. We refer to  $\varepsilon$  as the exploration rate. The value of  $\varepsilon_t$  in period  $t$  is given by

$$\varepsilon_t = 1 - e^{-\beta t} \quad (7)$$

where  $\beta$  is the convergence parameter that governs how quickly the algorithms shift from exploration to exploitation. As shown in Calvano et al. (2020b), the probability of collusion is generally increasing in  $\beta$ . The more the algorithms are allowed to explore, the more likely it is that they “stumble upon” a reward-punishment scheme. Once they discover these schemes, they are likely to adopt them since they are more profitable than playing Nash Equilibrium in the long run.

*Q Matrix.* Each algorithm policy function is defined by its own Q-matrix. The Q-matrix has dimension  $m^k \times m$  where the first dimension indicates the state space, i.e., the sequence of own past actions, and the second dimension indicates the current action. We initialize the  $Q$  matrix to the sum of discounted value of future profits, given action  $a_i$  and averaging over the possible actions of the opponents. Therefore, the initial values do not depend on the state  $\mathbf{s}$ , but only on the action  $a$ :

$$Q_i^0(\mathbf{s}, a_i) = \frac{1}{|\mathcal{A}|} \sum_{\mathbf{a}_{-i} \in \mathcal{A}^{n-1}} \frac{\pi_i(a_i, \mathbf{a}_{-i})}{1 - \delta} \quad (8)$$

*Policy Update.* Irrespectively of whether an algorithm is exploring or exploiting, the  $Q$  matrix is updated by averaging it's new value with the previous one, according to a parameter  $\alpha$ , the learning rate. In particular, the updating formula is the following:

$$Q_i(\mathbf{s}, a_i^*) = \alpha Q_i(\mathbf{s}, a_i^*) + (1 - \alpha) \left[ \pi(\mathbf{s}, a_i^*) + \delta \max_{a'_i} Q_i(\mathbf{s}', a'_i) \right]. \quad (9)$$

The new value of  $Q$  is state  $\mathbf{s}$  for action  $a_i^*$  is an average of the old value,  $Q_i(\mathbf{s}, a_i^*)$ , and the new one. The new value is the static payoff of action  $a_i^*$  in state  $\mathbf{s}$ , plus the discounted value of the best action the next state,  $\mathbf{s}'$ .

*Algorithm.* I summarize the full algorithm in Figure 1

---

**Algorithm 1:** Q-learning

---

```

initialize  $Q_i^0(\mathbf{s}, a_i) \forall i = 1 \dots n, \mathbf{s} \in \mathcal{S}, a_i \in \mathcal{A}$  ;
initialize  $\mathbf{s}^0$  ;
while convergence condition not met do
    exploration $_i = I(r_i < e^{-\beta t})$  where  $r_i \sim U(0, 1) \forall i$  ;
    if exploration $_i$  then
        |  $a_i^* = a \in A$  chosen uniformly at random ;
    else
        |  $a_i^* = \arg \max_{a_i} Q_i(\mathbf{s}, a_i)$  ;
    end
    determine  $\mathbf{s}'$  given  $(\mathbf{s}, a^*)$  ;
     $Q_i(\mathbf{s}, a_i^*) = \alpha Q_i(\mathbf{s}, a_i^*) + (1 - \alpha) [\pi(\mathbf{s}, a^*) + \delta \max_{a'_i} Q_i(\mathbf{s}', a'_i)] \forall i$  ;
     $\mathbf{s} = \mathbf{s}'$  ;
end

```

---

*Parametrization.* I summarize the parameters of the baseline model in Table 1. The parametrization closely follows Calvano et al. (2020b) so that the simulation results are directly comparable with theirs.

Parameter Description	Parameter	Value
Learning rate	$\alpha$	0.1
Exploration rate	$\beta$	$10^{-5}$
Discount factor	$\delta$	0.95
Marginal cost	$c$	1
Number of past observed states	$k$	1
Dimension of the action grid	$m$	5
Number of firms	$n$	2
Convergence parameter	$T$	$U[0, 1]$
Price elasticity	$\mu$	0.25

Table 1: Model parametrization

### 3.2 Convergence

There is no guarantee of convergence for the learning algorithm. Algorithms react to each others' policies and therefore it is possible that they get stuck in a loop. Moreover, as long as there is a non-zero probability of exploration, there is always a change that one algorithm embraces a totally different policy.

I use the same convergence criterion of Calvano et al. (2020b): convergence in actions. The algorithm stops if for  $T$  periods the index of the highest value of the  $Q$  matrix in each state has not changed i.e.  $\arg \max_{a_i} Q_{i,t}(\mathbf{s}_t, a_i) = \arg \max_{a_i} Q_{i,t+\tau}(\mathbf{s}_{t+\tau}, a_i) \forall i, \mathbf{s}, \tau = 1 \dots T$ . In practice, I choose a value of  $T = 10^5$ .

The advantage of this convergence criterion is that it does not require the algorithms to adopt a full exploitative strategy in order to converge. Algorithms' actions can stabilize even if the algorithms regularly take a random action with low probability. In fact, as long as the learning-rate  $\alpha$  is different from 1, firms only partially update their  $Q$  function. In practice, we observe convergence around 500,000 to 800,000 iterations, i.e. when the exploration probability  $\varepsilon$  is around  $e^{-5}$  to  $e^{-8}$ .

It is important to remark that one can achieve convergence also in a shorter amount of periods. What is crucial to achieve convergence is a sufficiently little exploration rate. However, with less periods, the algorithms are less likely to discover reward-punishment collusive strategies.

### 3.3 Results

From simulation results we see that the algorithms converge to supra-competitive prices. In Figure 1, I plot the distribution of equilibrium prices over 100 simulations. As we can see, the distribution is closer to monopoly prices than to the Nash Equilibrium prices. This also corresponds to higher profits.




Figure 1: Distribution of profits

These supra-competitive profits are achieved thanks to a reward-punishment scheme. Algorithms set higher prices until they observe a deviation. Once they observe a deviation, a punishment scheme starts and they earn lower profits for a couple of periods. Afterwards, they go back to the supracompetitive prices.

In order to see the reward-punishment scheme, I take the equilibrium  $Q$  function and manually manipulate the pricing action of one firm in one period and observe the reaction of both firms in the following periods. In particular, I make Algorithm 1 take the static best reply to Algorithm 2's action in the previous period. Then I let the two algorithms react to this unilateral deviation in the following periods. I report the sequence of prices of the two algorithms in Figure 2.




Figure 2: Actions when Algorithm 1 deviates

As we can see, when Algorithm 1 deviates from the stable collusive play, it sets a lower price. Since both algorithms observe only their own actions, Algorithm 2 does not deviate and keeps its collusive price. However, Algorithm 1 reacts to its own deviation and, in the following periods, it sets an even lower price before getting back to the stable collusive play.

In order to verify whether this was a reward-punishment scheme, we would need to observe that the profits of Algorithm 1 have increased in the deviation period, while they have decreased in the subsequent punishment periods. In Figure 3, I plot the profits of the two algorithms.






Figure 3: Profits when Algorithm 1 deviates

Figure 3 confirms our hypothesis of collusive play: the deviating algorithm achieves higher static profits in the deviation period but it is punished in the following periods.

In order to assess the robustness of this result, I report the time series of the average profits of the deviating algorithm over 100 simulations in Figure 4. The vertical bars represent interquartile ranges while the vertical lines represent minimum and maximum values.




Figure 4: Average profits when Algorithm 1 deviates

As we can see, Algorithm 1 consistently achieves higher profits in the deviating period, but is punished afterwards.

### 3.4 Comments

In the previous paragraph, we have seen how firms can learn to play reward-punishment schemes with the intent to keep supra-competitive prices even when their strategy is independent from the opponent action. Whether this behavior is collusive, it is subject to debate. Harrington (2018) tackles the issue of defining collusion in the presence of algorithmic pricing and lays down the following definition.

**“Definition:** Collusion is when firms use strategies that embody a reward–punishment scheme which rewards a firm for abiding by the supracompetitive outcome and punishes it for departing from it.”

According to this definition, the behavior we observed in the previous section, would be defined as collusive. What is puzzling is the fact that the strategy of the firm is independent of the opponent strategy and the punishment comes from the algorithm itself and not from its competitor. We conclude this section by examining how the firm learns these strategies and why they are stable.

First of all, how do the algorithms learn these reward-punishment schemes? As we have already seen in Section 2, reinforcement learning algorithms can potentially learn any strategy that is allowed by their state-action space. This means that their limits reside on what they can observe (the state space) and what they

can do (the action space). In this case, the algorithms observe only their own past action, their own price in the previous period, and use this information to set the current price. Therefore, they can learn any strategy in which the current price depends on the last price. The more they explore the state-action space, the more likely they are to learn strategies that bring higher payoffs.

Second, why are these strategies stable and why firms do not learn to undercut their rivals? The fact that algorithms are bounded by their state-action space is both a constraint and an opportunity. In fact, in the exploration phase, the algorithm can potentially learn any strategy compatible with its state-action space. However, when the algorithm shifts towards exploitation, the probability of exploring for multiple periods, decreases quickly. This means that the algorithm tests its current strategy under against and shorter deviations, in terms of how many sequential periods are explored. At a certain point, the algorithm will test its own current strategy only against one-shot deviations.

Let us explore in detail the case in which Algorithm 1 explores the possibility of best-replying to the static price of its rival. This happens with probability  $\varepsilon$ . We assume we are in “exploitation-most” mode so that  $\varepsilon$  is small and hence  $\varepsilon^2 \approx 0$ .

1. Algorithm 1 best replies to the supracompetitive price of Algorithm 2 setting  $p_1^{BR}(p_2^C)$
2. It gets a higher static payoff

$$\pi_1(p_1^{BR}(p_2^C), p_2^C) > \pi_1(p_1^C, p_2^C) \quad (10)$$

where  $p_1^C$  and  $p_2^C$  are the collusive prices of the two algorithms and  $p_1^{BR}(p_2^C)$  is the best reply of Algorithm 1 to the collusive price of Algorithm 2.

3. Algorithm 1 updates its  $Q$  function according to Equation 9:

$$Q_1(\mathbf{p}^C, p_1^{BR}(p_2^C)) = \alpha Q_1(\mathbf{p}^C, p_1^{BR}(p_2^C)) + (1 - \alpha) \left[ \pi_1(p_1^{BR}(p_2^C), p_2^C) + \delta \max_{p'_1} Q_1(p_1^C, p_1^{BR}(p_2^C), p'_1) \right] \quad (11)$$

where  $\mathbf{p}^C$  is the vector of collusive prices.

4. As we can see, the update depends of two terms: the static profits  $\pi_1(p_1^{BR}(p_2^C), p_2^C)$  and the future value  $\delta \max_{p'_1} Q_1(p_1^C, p_1^{BR}(p_2^C), p'_1)$ . Even if static profits are bigger than current profits, the future value is not. Why? Because of the punishment scheme. Algorithm 1 will choose  $p'_1 = p_1^P$  so that it will punish itself and, as a consequence, it will not see  $p_1^{BR}(p_2^C)$  as a better strategy.
5. This would be different if Algorithm 1, in the next period would explore again and pick  $p_1^{BR}(p_2^C)$  instead of  $p_1^P$ . However, as we have said at the beginning, this is extremely unlikely since we are in “exploitation-most” mode so that  $\varepsilon^2 \approx 0$ .

Third, how do algorithms converge to these strategies? My claim is that the crucial ingredient so that algorithms simultaneously converge to collusive strategies is synchronous learning. If the algorithms were learning asynchronously, i.e., one algorithm is in “exploitation-most” mode while the other is in “exploration-most” mode, we would not obtain the same result. The algorithm in “exploration-most” mode would learn a new strategy to exploit the collusive algorithm. I test this hypothesis in the next section.

## 4 Detecting Algorithmic Collusion

In the previous section, we have seen that detecting algorithmic collusion from the inspection of the algorithm’s inputs might not be feasible. Even in a very simple setting, where inspection of the algorithm’s inputs seems

sufficient to determine whether they can learn reward-punishment schemes with the intent to keep supracompetitive prices, simple rules might not be sufficient. Algorithms learn to collude even when they cannot base their strategy on past rival’s actions.

One solution could be to detect collusion from observational data. Economists know signs that could hint at collusive behavior such as coordinated prices. However, these methods rely heavily on the underlying models of market interactions. In fact, one has to be able to distinguish collusion from a wide variety of confounding factors such as demand, aggregate shocks, input prices or simply noise.

In this section, I study a model-free method to detect algorithmic collusion. My method is based on the observation that the crucial aspect underlying algorithmic collusion is synchronous learning. I test the hypothesis that algorithms can learn to unilaterally exploit collusive behavior, but they need to be retrained from scratch in order to learn these strategies. I show that unilaterally retraining one algorithm, after they have learned collusive strategies, leads to more competitive prices. I show this result both in the original setting of Calvano et al. (2020b) and in the setting of Section 3.

#### 4.1 Blind Collusive Algorithm Retraining

If firms are able to learn complex grim-trigger strategies, they should also be able to learn to exploit another algorithm that is playing a collusive strategy. In this section, I explore this possibility. In particular, I explore what happens when one of the two algorithms gets unilaterally re-trained.

I first train the two algorithms as explained in Section 3. Once convergence is achieved, I manually restart the timing  $t$  for one of the two algorithms. This means that, according to Equation 7, the probability of exploration is reset to one and it will gradually shift towards zero as  $t$  increases. The algorithm will move from an “exploitation mode” mode to an “exploration most” mode. I do not reset the timing of the other algorithm which will proceed onwards, starting from the time when it converged. I use the same criterion described in Section 3.2 to determine convergence.

In Figure 5, I plot the actions of Algorithm 1 before and after re-training.




Figure 5: Average actions of Algorithm 1 with retraining

From the figure, we can see that indeed the algorithm learns a different, more competitive strategy. In Figure 6, I plot the firms’ profits before and after re-training.




Figure 6: Average profits of Algorithm 1 with retraining

As we can see, retraining leads to a more profitable strategy for Algorithm 1. The algorithm learns to exploit the collusive strategy of its opponent and undercuts it, earning higher profits. The opponent does not observe Algorithm 1's action and therefore sticks to its own current strategy, as we can see from Figure 7.

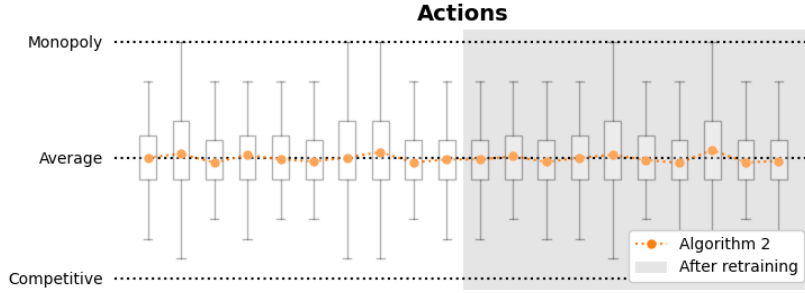


Figure 7: Average actions of Algorithm 2 with retraining of Algorithm 1

## 4.2 Non-blind Collusive Algorithm Retraining

In this section, I show that the same happens in the original framework of Calvano et al. (2020b), when the algorithms do not only observe their own past action but also those of their opponent. In this case, the algorithms can actively enforce punishment in case of rival's deviation. As a consequence, the best strategy in response to the collusive policy might not be as simple as undercutting. In fact, the collusive policy includes a punishment action for undercutting. I plot the actions of Algorithm 1 before and after retraining in Figure 8.




Figure 8: Average actions of Algorithm 1 with retraining

In Figure 8, we observe that Algorithm 1 learns a complex dynamic strategy in order to exploit the collusive policy of Algorithm 2. Importantly, this strategy is more competitive than the collusive one. In order to see whether this strategy is indeed an improvement over the collusive one, I plot the profits of Algorithm 1 before and after retraining in Figure 9.




Figure 9: Average profits of Algorithm 1 with retraining

Figure 9 confirms that the new policy of Algorithm 1 after retraining is indeed more profitable than the collusive policy. If now Algorithm 2 was to be retrained, we would see it adopting an even more competitive strategy. Iterating retraining over and over, the two algorithms would converge to Nash Equilibrium prices.

### 4.3 Non Collusive Algorithm Retraining

Lastly, we are going to inspect what happens when one algorithm is unilaterally retrained in case the algorithms were not colluding. We would expect that if the algorithms were play competitive strategies, one algorithm would not be able to learn any better strategy than the competitive one. In Figure 10, I plot the actions before and after retraining of Algorithm 1, in the scenario in which the algorithms were playing competitive strategies.




Figure 10: Average profits of Algorithm 1 with retraining

As we can see from Figure 10, Algorithm 1 keeps playing the competitive strategy even after re-training.

### 4.4 Re-training on historical data

As we have seen in the previous paragraphs, asynchronous learning leads to more competitive strategies. However, the direct policy implication of this finding - enforcing retraining of algorithms - is clearly unfeasible. One cannot ask firms to reset their algorithms in order to check for collusion. Learning is not only time consuming but also expensive for firms that have to undergo a sequence of suboptimal actions.

In this section, I show that one can obtain similar insights also from observational data. The key insight comes from the fact that algorithm exploration provides natural experiments to retrain the algorithm. In fact, thanks to algorithm exploration, firms collect counterfactual data on many possible scenarios. Given this data, one could retrain the algorithm from scratch. In particular, in order to perform retraining, one needs the following ingredients: for each state-action pair  $(s, a_i)$ , a list of payoffs  $\Pi_i(s, a_i)$  and future states  $S'(s, a_i)$ . With these ingredients, one can run the algorithm depicted in Figure 1, where payoffs and next states are drawn uniformly at random from  $\Pi_i(s, a_i)$  and  $S'(s, a_i)$  respectively. I sketch the structure of the

bootstrapping algorithm in Figure 2.

---

**Algorithm 2:** Q-learning from observational data

---

```

initialize  $Q_i^0(\mathbf{s}, a_i) \forall i = 1 \dots n, \mathbf{s} \in \mathcal{S}, a_i \in \mathcal{A}$  ;
initialize  $\mathbf{s}^0$  ;
while convergence condition not met do
    exploration $i$  =  $I(r_i < e^{-\beta t})$  where  $r_i \sim U(0, 1)$  ;
    if exploration $i$  then
        |  $a_i^* = a \in \mathcal{A}$  chosen uniformly at random ;
    else
        |  $a_i^* = \arg \max_{a_i} Q_i(\mathbf{s}, a_i)$  ;
    end
    uniformly at random draw a period  $o$  from the observed periods ;
    select  $\mathbf{s}'(\mathbf{s}, a_i) = \mathbf{S}'(\mathbf{s}, a_i, o)$  ;
     $Q_i(\mathbf{s}, a_i^*) = \alpha Q_i(\mathbf{s}, a_i^*) + (1 - \alpha) [\Pi_i(\mathbf{s}, a_i^*, o) + \delta \max_{a'_i} Q_i(\mathbf{s}', a'_i)]$  ;
     $\mathbf{s} = \mathbf{s}'$  ;
end

```

---

The choice of the periods is important. The more recent one observation is, the more likely it is to be representative of current opponent's behavior. However, one would also prefer to include in the bootstrap sample older observations just to increase the number of observations. In the simulation results reported below, for each state-action pair  $(\mathbf{s}, a_i)$ , I use the 3 most recent period to retrain the algorithm.

It is important to note that, given the choice of the bootstrap sample, the optimization problem becomes stationary. Therefore, in this case, one can actually apply value function iteration in order to find the value of  $Q$ . Moreover, the process is now guaranteed to converge.

In Figure 11, I report the actions of Algorithm 1 after retraining of the bootstrap sample, in the setting of Section 3, when the algorithm sees only its own past action.




Figure 11: Average actions of Algorithm 1 with retraining from bootstrap data

From Figure 11, we observe that the algorithm consistently changes its strategy adopting more competitive actions. The results are more noisier than in the case of retraining, but this is expected since observed actions are only partially predictive of future behavior. In Figure 12, I repeat the same exercise in the setting of Calvano et al. (2020b), when the algorithms observe both their own past action and their competitor's.




Figure 12: Average actions of Algorithm 1 with retraining from bootstrap data

Again, Algorithm 1 consistently picks more competitive actions after retraining from the bootstrap sample of observed data.

## 5 Conclusion

In this paper, I have examined the issue of how to detect algorithmic collusion. I have first shown that the inspection of the algorithm’s inputs might not be sufficient in order to determine whether an algorithm can learn collusive strategies. Using the algorithmic collusion setting of Calvano et al. (2020b), I show that even if algorithms were not observing the competitor’s actions, they are still able to learn reward-punishment schemes. While this is an extreme scenario, it highlights the fact that independence from competitors’ actions is not sufficient to prevent algorithms to learn reward-punishment strategies with the purpose to set supra-competitive prices.

In the second part of the paper, I propose a model-free approach to detect algorithmic collusion. To the best of my knowledge, this is the first attempt to build a model-free test for algorithmic collusion, using only historical data. Building on the insight that the key aspect element that allows algorithms to learn to collude is synchronous learning, I show that unilaterally re-training of one algorithm leads to more competitive strategies. Remarkably, this happens only if the algorithms were colluding. The same does not occur if the algorithms were adopting competitive strategies.

Since retraining algorithms is expensive and difficult to enforce, I show that one can obtain similar results relying only on historical data. In fact, a key feature of reinforcement learning is that the algorithm keeps exploring new strategies in order to adapt to changes in a dynamic environment. These exploration phases provide natural experiments to test ex-post for collusive behavior.

## References

- Stephanie Assad, Robert Clark, Daniel Ershov, and Lei Xu. Algorithmic pricing and competition: Empirical evidence from the german retail gasoline market. 2020.
- Alan W Beggs. On the convergence of reinforcement learning. *Journal of economic theory*, 122(1):1–36, 2005.
- Emilio Calvano, Giacomo Calzolari, Vincenzo Denicolò, Joseph E Harrington, and Sergio Pastorello. Protecting consumers from collusive prices due to ai. *Science*, 370(6520):1040–1042, 2020a.
- Emilio Calvano, Giacomo Calzolari, Vincenzo Denicolo, and Sergio Pastorello. Artificial intelligence, algorithmic pricing, and collusion. *American Economic Review*, 110(10):3267–97, 2020b.
- Le Chen, Alan Mislove, and Christo Wilson. An empirical analysis of algorithmic pricing on amazon marketplace. In *Proceedings of the 25th International Conference on World Wide Web*, pages 1339–1349, 2016.

- Ariel Ezrachi and Maurice E Stucke. Artificial intelligence & collusion: When computers inhibit competition. *U. Ill. L. Rev.*, page 1775, 2017.
- Ariel Ezrachi and Maurice E Stucke. Sustainable and unchallenged algorithmic tacit collusion. *Nw. J. Tech. & Intell. Prop.*, 17:217, 2019.
- Joseph E Harrington. Developing competition law for collusion by autonomous artificial agents. *Journal of Competition Law & Economics*, 14(3):331–363, 2018.
- Mitsuru Igami. Artificial intelligence as structural estimation: Deep blue, bonanza, and alphago. *The Econometrics Journal*, 23(3):S1–S24, 2020.
- Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.